



Database Management Systems - I, CS 157A

SQL Overview and SELECT



Agenda

- Create Schema (DDL)
- Query
 - Select-From-Where Statements
 - Multi-relation Queries
 - Sub-queries

Review: A Relation is a Table

Attributes
(column
headers)

Tuples
(rows)

name	manf
Winterbrew	Pete's
Bud Lite	Anheuser-Busch

Relation
name

Beers



Review: From E/R Diagrams to Relations

- **Entity set** → relation.
 - Attributes → attributes.
- **Relationships** → relations whose attributes are only:
 - The keys of the connected entity sets.
 - Attributes of the relationship itself.



Review: Core Relational Algebra

- **Union, intersection, and difference:**
 - Usual set operations, but *both operands must have the same relation schema*.
- **Selection:** picking certain rows.
- **Projection:** picking certain columns.
- **Products** and **Joins:** compositions of relations.
- **Renaming** of relations and attributes.



Define Schemas - DDL



Database Schemas in SQL

- Structured Query Language (SQL)
- SQL is primarily a *query* language, for getting information from a database.
- But SQL also includes a *data-definition* component for describing database schemas.



SQL Statements

DML (Data Manipulation Language)	SELECT
	INSERT UPDATE DELETE
DDL (Data Definition Language)	CREATE ALTER DROP
DCL and Transaction Control	GRANT REVOKE COMMIT ROLLBACK



Creating (Declaring) a Relation

- Simplest form is:

```
CREATE TABLE <name> (  
    <list of elements>  
);
```

- To delete a relation:

```
DROP TABLE <name>;
```



Example: Create Table

```
CREATE TABLE Sells (  
    bar      CHAR(20),  
    beer     VARCHAR(20),  
    price    REAL  
);
```



Declaring Single-Attribute Keys

- Place PRIMARY KEY or UNIQUE after the type in the declaration of the attribute.
- Example:

```
CREATE TABLE Beers (  
    name        CHAR(20) UNIQUE,  
    manf        CHAR(20)  
);
```



Example: Multi-attribute Key


- The bar and beer together are the key for *Sells*:

```
CREATE TABLE Sells (  
    bar          CHAR(20) ,  
    beer         VARCHAR(20) ,  
    price        REAL ,  
    PRIMARY KEY (bar, beer)  
);
```



PRIMARY KEY vs. UNIQUE

1. There can be only one PRIMARY KEY for a relation, but several UNIQUE attributes.
2. No attribute of a PRIMARY KEY can ever be NULL in any tuple. But attributes declared UNIQUE may have NULL's, and there may be several tuples with NULL.



Query



SQL Query

- SQL is a very-high-level language:
 - Say “**what to do**” rather than “how to do it.”
 - Avoid a lot of data-manipulation details needed in procedural languages like C++ or Java.
- Database management system figures out “best” way to execute query:
 - Called “query optimization.”



Select-From-Where Statements

SELECT desired attributes

FROM one or more tables

WHERE condition about tuples of
the tables



Our Running Example

- All our SQL queries will be based on the following database schema.
 - Underline indicates key attributes.

Beers(name, manf)

Bars(name, addr, license)

Drinkers(name, addr, phone)

Likes(drinker, beer)

Sells(bar, beer, price)



Example

- Using **Beers(name, manf)**, what beers are made by Anheuser-Busch?

```
SELECT name
FROM   Beers
WHERE  manf = 'Anheuser-Busch' ;
```



Result of Query

name
Bud
Bud Lite
Michelob
. . .

The answer is a relation with a single attribute, name, and tuples with the name of each beer by Anheuser-Busch, such as Bud.



Meaning of Single-Relation Query

- Begin with the relation in the **FROM** clause.
- Apply the selection indicated by the **WHERE** clause.
- Apply the extended projection indicated by the **SELECT** clause.

Operational Semantics

name	manf
Bud	Anheuser-Busch

Tuple-variable t
loops over all
tuples

Include $t.name$
in the result, if so

Check if
Anheuser-Busch



“*” In SELECT clauses

- When there is one relation in the FROM clause, * in the SELECT clause stands for “all attributes of this relation.”
- Example: Using Beers(name, manf):

```
SELECT *
```

```
FROM Beers
```

```
WHERE manf = 'Anheuser-Busch' ;
```



Result of Query:

name	manf
Bud	Anheuser-Busch
Bud Lite	Anheuser-Busch
Michelob	Anheuser-Busch
.

Now, the result has each of the attributes of Beers.



Renaming Attributes

- If you want the result to have different attribute names, use “**AS <new name>**” to rename an attribute.
- **Example:** Using **Beers(name, manf):**

```
SELECT name AS beer, manf
FROM    Beers
WHERE   manf = 'Anheuser-Busch' ;
```




Result of Query:

beer	manf
Bud	Anheuser-Busch
Bud Lite	Anheuser-Busch
Michelob	Anheuser-Busch
.



Expressions in SELECT Clauses

- Any expression that makes sense can appear as an element of a SELECT clause.
- **Example:** Using **Sells(bar, beer, price):**

```
SELECT bar, beer,  
       price*114 AS priceInYen  
FROM   Sells;
```



Result of Query

bar	beer	priceInYen
Joe's	Bud	285
Sue's	Miller	342
...

Example: Constants as Expressions

- Using Likes(drinker, beer):

```
SELECT drinker,      Attribute value      Attribute name
           'likes Bud' AS whoLikesBud
FROM Likes
WHERE beer = 'Bud' ;
```



Result of Query

drinker	whoLikesBud
Sally	likes Bud
Fred	likes Bud
...	...



Example: Information Integration

- We often build “data warehouses” from the data at many “sources.”
- Suppose each bar has its own Menu relation **Menu(beer, price)** .
- To contribute to **Sells(bar, beer, price)** we need to query each bar and insert the name of the bar.



Information Integration (cont.)

- For instance, at Joe's Bar we can issue the query:

```
SELECT  'Joe''s Bar', beer,  
        price  
FROM    Menu;
```



Complex Conditions in WHERE Clause

- Boolean operators AND, OR, NOT.
- Comparisons =, <>, <, >, <=, >=.
- And many other operators that produce boolean-valued results.



Example: Complex Condition

- Using **Sells(bar, beer, price)**, find the price Joe's Bar charges for Bud:

```
SELECT price
FROM Sells
WHERE bar = 'Joe's Bar' AND
      beer = 'Bud';
```



Patterns

- A condition can compare a string to a pattern by:
 - `<Attribute> LIKE <pattern>` or `<Attribute> NOT LIKE <pattern>`
- *Pattern* is a quoted string with
 - `%` = “any string.”
 - `_` = “any single character.”



Example: LIKE

- Using **Drinkers(name, addr, phone)** find the drinkers with exchange 555:

```
SELECT name
FROM   Drinkers
WHERE  phone LIKE '%555-__ __ __';
```



NULL Values

- Tuples in SQL relations can have NULL as a value for one or more components.
- **Meaning depends on context. Two common cases:**
 - **Missing value:** e.g., we know Joe's Bar has some address, but we don't know what it is.
 - **Inapplicable:** e.g., the value of attribute **spouse** for an unmarried person.



Comparing NULL's to Values

- The logic of conditions in SQL is really 3-valued logic: **TRUE**, **FALSE**, **UNKNOWN**.
- Comparing any value (including **NULL** itself) with **NULL** yields **UNKNOWN**.
- A tuple is in a query answer iff the **WHERE** clause is **TRUE** (not **FALSE** or **UNKNOWN**).

Surprising Example


- From the following Sells relation:

bar	beer	price
Joe's Bar	Bud	NULL

SELECT bar

FROM Sells

WHERE price < 2.00 OR price >= 2.00;





Multi-table Queries

- Interesting queries often combine data from more than one table.
- We can address several tables in one query by listing them all in the **FROM** clause.
- Distinguish attributes of the same name in 2 tables by using “<table>.<attribute>” .



Example: Joining Two tables

- Using tables **Emp(ename, dno)** and **Dept(dno, dname)**, find the department name of employee *Joe*.


```
SELECT dname
FROM    Emp, Dept
WHERE   ename = 'Joe' AND
        Emp.dno = Dept.dno;
```




Operational Semantics

- Imagine one tuple-variable for each relation in the **FROM** clause.
 - These tuple-variables visit each combination of tuples, one from each relation.
- If the tuple-variables are pointing to tuples that satisfy the **WHERE** clause, send these tuples to the **SELECT** clause.

Example





Example: Self-Join

- From **Beers(name, manf)**, find all pairs of beers by the same manufacturer.
 - Do not produce pairs like (Bud, Bud).
 - Produce pairs in alphabetic order, e.g. (Bud, Miller), not (Miller, Bud).

```
SELECT b1.name, b2.name
FROM   Beers b1, Beers b2
WHERE  b1.manf = b2.manf AND
       b1.name < b2.name;
```



SQL Sub-query



Sub-queries

- A parenthesized **SELECT-FROM-WHERE** statement (*subquery*) can be used as a value in a number of places, including FROM and WHERE clauses.
- **Example:** in place of a table in the WHERE clause, we can use a subquery and then query its result.

Example: Subquery in WHERE

- Find the department name that has at least one employee.

Emp(ename, dno) and Dept(dno, dname) relations

SELECT dname

FROM Dept

WHERE dno **IN**

(select dno
from Emp)

All employee's dno



Subqueries That Return One Tuple

- If a subquery is guaranteed to produce one tuple, then the subquery can be used as a value:
 - Usually, the tuple has one component.
 - A run-time error occurs if there is no tuple or more than one tuple.



Example: Single-Tuple Subquery

- Using tables **Emp(ename, dno, sal)** and **Dept(dno, dname)**, find the name of employee who gets the highest salary

```
SELECT ename
```

```
FROM Emp
```

```
WHERE sal >= (select max(sal)  
              from Emp);
```




The IN Operator

- **<tuple> IN (<subquery>)** is true if and only if the **tuple** is a member of the table produced by the subquery.
 - **Opposite:** **<tuple> NOT IN (<subquery>)**.
- IN-expressions can appear in **WHERE** clauses.

IN is a Predicate About R's Tuples

SELECT a

FROM R

WHERE b IN

(SELECT b FROM S);

Two 7's

One loop, over
the tuples of R

a	b
1	7
3	4

R

b	c
7	5
7	6

S

(1,7) satisfies
the condition;
1 is output once.

This Query Pairs Tuples from R, S

```
SELECT a
FROM   R, S
WHERE  R.b = S.b;
```

a	b
1	7
3	4

R

b	c
7	5
7	6

S

Double loop, over
the tuples of R and S

(1,7) with (7,5)
and (1,7) with
(7,6) both satisfy
the condition;
1 is output twice.



The Exists Operator

- **EXISTS**(<subquery>) is true if and only if the subquery result is not empty.
- **Example**: From **Emp(ename, dno)** , find those employees that are the only employee in their department.

Example: (Not) EXISTS

```
SELECT  ename
FROM    Emp e1
WHERE  NOT EXISTS (
    SELECT  *
    FROM    Emp
    WHERE    dno = e1.dno  AND
             ename <> e1.ename);
```

Notice scope rule: dno refers to closest nested FROM with a table having that attribute.

Notice the SQL "not equals" operator



Union, Intersection, and Difference

- Union, intersection, and difference of tables are expressed by the following forms, each involving subqueries:
 - (<subquery> **UNION** (<subquery>)
 - (<subquery> **INTERSECT** (<subquery>)
 - (<subquery> **EXCEPT** (<subquery>))



DISTINCT

- From `Sells(bar, beer, price)`, find all the different prices charged for beers:


```
SELECT DISTINCT price
FROM Sells;
```

- Notice that without `DISTINCT`, each price would be listed as many times as there were bar/beer pairs at that price.



Summary

- DDL - Create Schema
- Query
 - Select-From-Where Statements
 - Multi-relation Queries
 - Sub-queries



END